
ECE 391 Final Review

You're almost done, hang in there!

Reminders

- Final is Tuesday Dec 12, 7-10 PM
 - 3 sheets of notes
 - Topics Covered
 - Signals
 - I-Mail
 - Memory Allocation
 - Scheduling
 - Memory Maps
 - MP3
 - Final is not cumulative
-

Signals

What are Signals?

- User “Interrupts”
 - Signals can be generated by user processes or kernel
 - Signals can be handled either by kernel or user processes
 - Some signals can be masked (Non maskable signals exist and you should know them)
 - Signals vs Interrupts:
 - When do we check for signals – returning from:
 - Interrupt, Exception, syscall
 - Who manages them: Signals -> Kernel, Interrupts -> Processor (IDT)
 - Who handles them: Signals -> Kernel/User, Interrupts -> Kernel
-

Important Signals

- SIGINT (Ctrl-C)
 - SIGSTP (Ctrl-Z)
 - SIGCONT
 - SIGKILL
 - SIGSTOP
-
- Which two signals cannot have their default behavior changed?
-

Default Actions

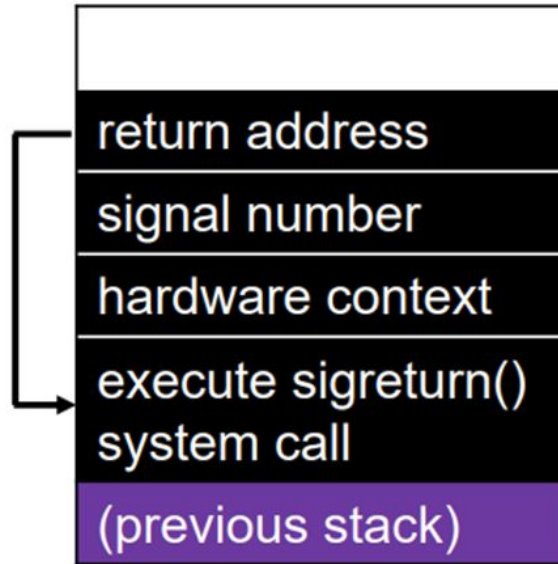
- Terminate
 - Ends the program completely
 - Dump Core
 - Terminates and dumps core
 - Stop
 - Stops a program, program can be resumed later
 - Ignore
 - Signal is discarded
 - Continue
 - Continues a currently stopped process
-

How does this sigreturn work?

- Default handler goes through kernel
 - If handler is assigned, kernel will “help” the user with returning to kernel
 - Building the user return argument
 - Must call another system call to get back to sigreturn
 - Kernel will set this up, rather than having user handler do the syscall
 - Swapping hardware context on return
 - Want: Get to signal code, Problem: sigreturn system call context not correct to get there
 - Solution: Store hardware context from first signal call interrupt, restore in sigreturn
-

Sigreturn User Stack

User stack after
signal is delivered



Things to Know

- Signal User stack
 - Possible coding question here
 - Who creates signals?
 - Signal control flow (Signal generator -> Kernel -> handler (user or kernel) -> sigreturn -> initial signal generator)
 - Non-Maskable signals
 - Signal related functions
-

I-Mail

What is a device driver?

- Kernel interacts with I/O devices through a device driver
 - Why?
 - Hides implementation details of how device works
 - Allows dynamic loading/unloading of device drivers
 - Creates a standardize API for interacting with the device
-

Examples

- Block devices:
 - Data is accessible in blocks of a fixed size (Few kB)
 - Data transfers are buffered and cached
 - Character devices:
 - Data is access at a byte level
-

Overview

- What is Blocking?
 - Waits until information is received to return
 - Functions that can block
 - Wait for a new message
 - Read
 - Poll
 - Write
 - Doesn't have to wait for device
 - Unique to I-mail
-

Overview (pt. 2)

- Sleeping
 - Step 1: Release All locks
 - Step 2: Ensure Conditions without locks are in secured state
 - Step 3: Go to Sleep
 - Step 4: Wake up
 - Reacquire locks
 - Check validity
-

Overview (pt. 3)

- Important Data Structures
 - Wait Queues
 - Efficient way to wait for previous tasks to finish
 - Formed via doubly linked list
 - Reduces CPU cycles
 - Could Create Race conditions
 - I-mail Read
 - Check for readable information & available semaphores
 - Read information as necessary
 - Release used sempahores
-

Dynamic Memory Allocation

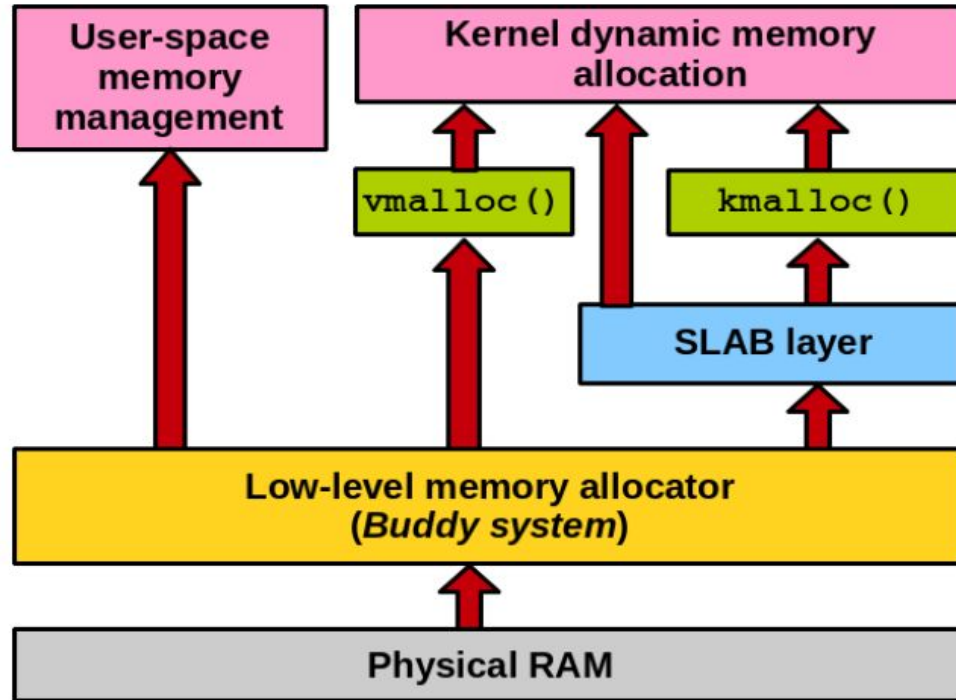
- User creates data in driver
 - Deleted by:
 - I-mail
 - Delete user function
 - Deleting user could create function if user is in the process of using I-mail
-

Synchronization

- Functions need to be synchronized appropriately
 - Think about which functions interact with each other
 - Think about order that operations will be conducted in
 - Use semaphores appropriately
-

Memory Allocation

Overview of allocation in Linux





Overview of allocation in Linux

- Physically contiguous:
 - Buddy allocator.
 - Slab allocator.
 - `kmalloc()`
- Non-physically contiguous (but virtually contiguous):
`vmalloc()`.


Buddy allocator interface

- No need to memorize function names, just know that buddy allocator serves up requested number of pages:
 - unsigned long **__get_free_page** (unsigned int gfp_mask)
 - Allocate a single page and return a virtual address
 - unsigned long **__get_free_pages** (unsigned int gfp_mask, unsigned int order)
 - Allocate 2^{order} number of pages and return a virtual address
-


Buddy allocator



Order 0 Pairs





Order 1 Pairs




Order 2 Pairs


Allocate Order 0 Block



Order 0 Pairs





Order 1 Pairs




Order 2 Pairs


Allocate Order 0 Block



Order 0 Pairs




Order 1 Pairs



Order 2 Pairs

Allocate Order 0 Block



Allocate Order 0 Block

Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block

0	0	0	0
---	---	---	---

Order 0 Pairs

0	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Allocate Order 0 Block

Allocated
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block
Order 0 Free Block

1	0	0	0
---	---	---	---

Order 0 Pairs

0	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Allocate Order 0 Block

Allocated
Order 0 Free Block
Order 1 Free Block
Order 1 Free Block
Order 1 Free Block

1	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Allocate Order 0 Block

Allocated
Order 0 Free Block
Order 1 Free Block
Order 2 Free Block

1	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

1

Order 2 Pairs

Allocate Another Order 0 Block

Allocated
Order 0 Free Block
Order 1 Free Block
Order 2 Free Block

1	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

1

Order 2 Pairs

Allocate Another Order 0 Block

Allocated
Allocated
Order 1 Free Block
Order 2 Free Block

0	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

1

Order 2 Pairs

Allocate Order 2 Block

Allocated
Allocated
Order 1 Free Block
Order 2 Free Block

0	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

1

Order 2 Pairs

Allocate Order 2 Block

Allocated
Allocated
Order 1 Free Block
Allocated

0	0	0	0
---	---	---	---

Order 0 Pairs

1	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Free Order 0 Block

Allocated
Allocated
Order 1 Free Block
Allocated

0	0	0	0
---	---	---	---

Order 0 Pairs



1	0
---	---

Order 1 Pairs


0

Order 2 Pairs


Free Order 0 Block



Order 0 Pairs



Order 1 Pairs



Order 2 Pairs

Free Other Order 0 Block

Allocated
Order 0 Free Block
Order 1 Free Block
Allocated

1	0	0	0
---	---	---	---

Order 0 Pairs



1	0
---	---

Order 1 Pairs


0

Order 2 Pairs


Free Other Order 0 Block



Order 0 Pairs





Order 1 Pairs




Order 2 Pairs


Free Other Order 0 Block



Order 0 Pairs





Order 1 Pairs




Order 2 Pairs


Free Other Order 0 Block



Order 0 Pairs





Order 1 Pairs




Order 2 Pairs


Free Order 2 Block



Order 0 Pairs



Order 1 Pairs



Order 2 Pairs

Free Order 2 Block

Order 2 Free Block

Order 2 Free Block

0	0	0	0
---	---	---	---

Order 0 Pairs

0	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Free Order 2 Block

Order 3 Free Block

0	0	0	0
---	---	---	---

Order 0 Pairs

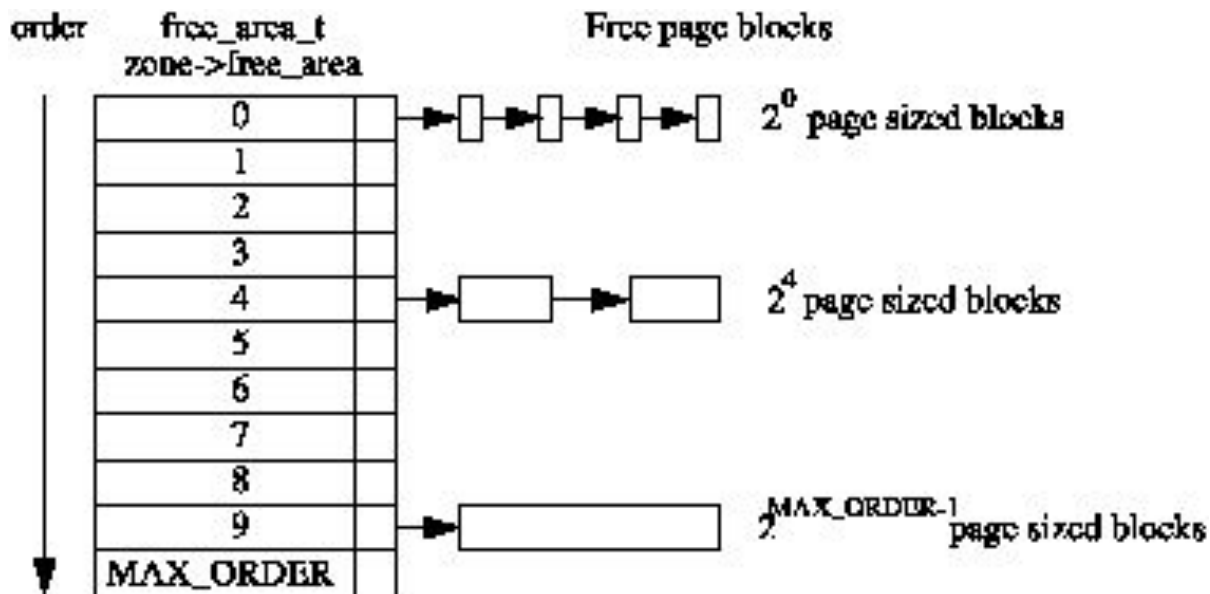
0	0
---	---

Order 1 Pairs

0

Order 2 Pairs

Free Page List



Buddy allocation issues

- Buddy allocation **reduces external fragmentation** (how?)
 - However, it **doesn't prevent internal fragmentation** (why?)
 - Request 33 pages, best block it'll give you is 64 pages. 31 pages = 48% wasted.
 - Linux's solution: layer "slab allocator" on top of buddy allocator to solve internal fragmentation.
-

Slab allocation: the big idea

- Some structs are allocated and freed very often.

Examples:

- task_struct, mutex's, inodes, dentries
 - There's an overhead to destructing and constructing these objects over and over.
 - Solution: **cache!** [Bonwick94]
-

Slab allocation: 2 principal goals

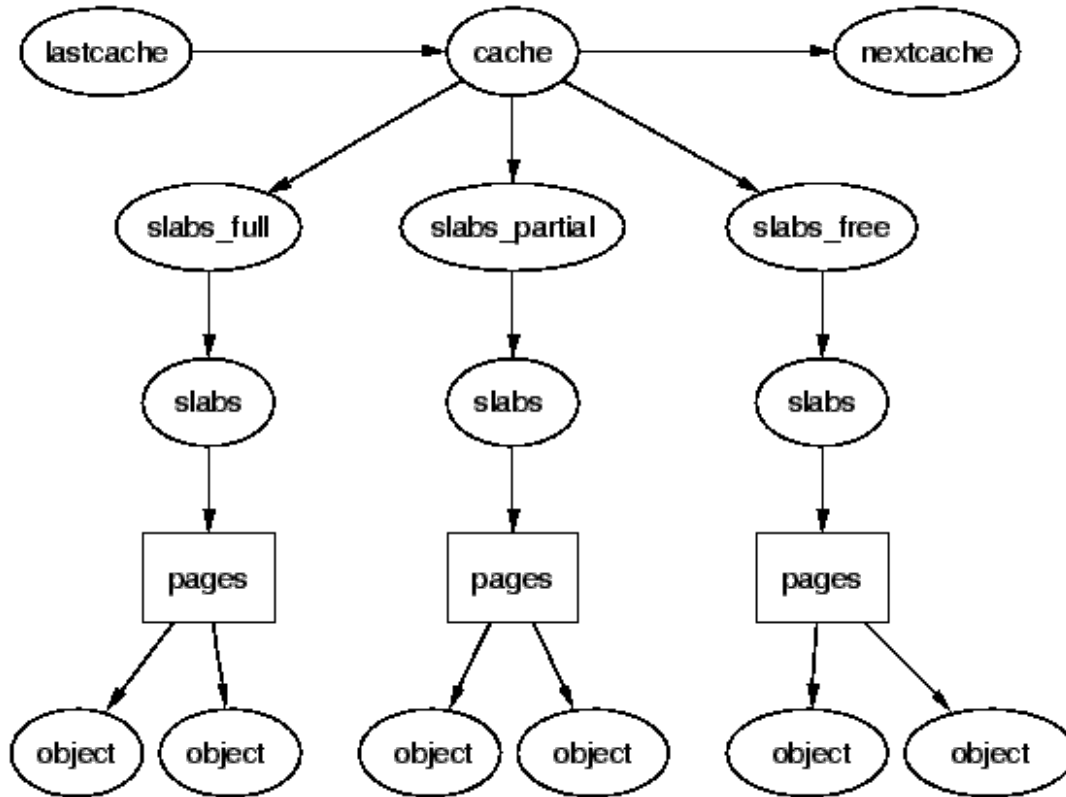
- Allow allocation of small blocks of memory to help eliminate internal fragmentation that would be otherwise caused by the buddy system.
 - Have caches of commonly used objects kept in an initialised state available for use by the kernel.
 - Additional principle: HW cache-align objects to speed up access.
-

Object caching: basic algorithm

```
// Allocate an object
if (there's an object in the cache)
    take it (no construction required);
else {
    allocate memory;
    construct the object;
}

// Free an object
return it to cache (no destruction --
                    simply return it
                    to init state);
```

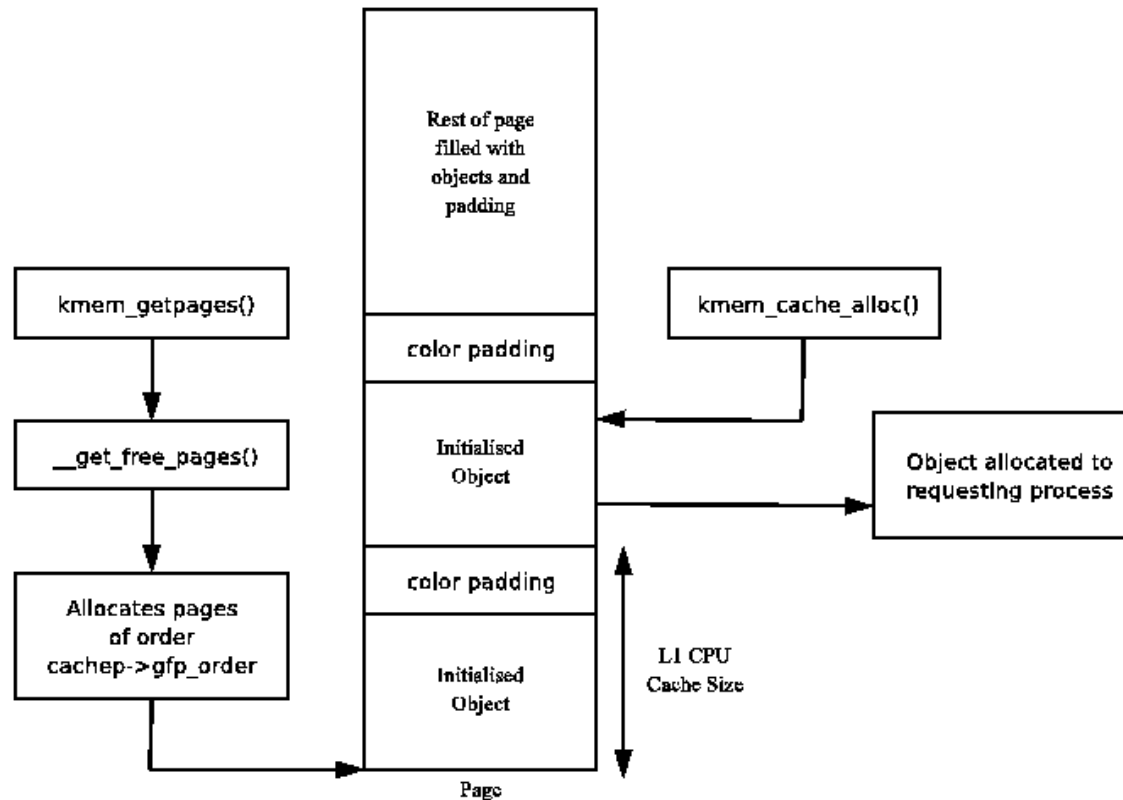
Slab cache structure: one cache per object type



Definitions:

- Slab: one or more pages of physically contiguous memory carved up into equal-size chunks.
- Slab cache: collection of slabs that correspond to a single object type.

Slab allocation: slightly more detail



Ex: your object type is a 24B struct. Your slab is 8192B (2 pages). Assuming no color padding, how many structs can you fit in the slab?

$\text{floor}(8192/24) = 341$.

kmalloc: lean on slab cache.

- Kernel has a separate bunch of caches specifically for kmalloc.
 - On a Linux box, try: `sudo vmstat -m | grep kmalloc.`
 - These range from ~8K to 8B. (depends on arch).
 - When kmalloc is called, it uses one of the dedicated slab caches to fulfill the request.
 - Leaning on slab caches helps minimize internal fragmentation within a page, and allows dynamic allocation of small buffers.
-

```
$ sudo vmstat -m | grep "kmalloc"
```

Cache	Num	Total	Size	Pages
...				
kmalloc-8k	297	312	8192	4
kmalloc-4k	1028	1072	4096	8
kmalloc-2k	1655	1664	2048	16
kmalloc-1k	2711	2720	1024	32
kmalloc-512	34017	34592	512	32
kmalloc-256	13146	13344	256	32
kmalloc-192	3309	3402	192	42
kmalloc-128	3609	3712	128	32
kmalloc-96	6621	7854	96	42
kmalloc-64	20124	21120	64	64
kmalloc-32	15363	15744	32	128
kmalloc-16	28787	31232	16	256
kmalloc-8	15323	15360	8	512
...				

vmalloc

- If we want lots of memory, finding physically contiguous memory may be tricky despite buddy allocator.
 - Linux provides vmalloc, which gets a large virtually contiguous chunk of memory but not necessarily contiguous physically.
 - **vmalloc/vfree** is the interface.
-

Scheduling

Why do we need scheduling?

- Scheduling provides the *illusion* of multiple processes running simultaneously without additional hardware requirements
 - Ensures fairness of CPU usage
 - Make computers more useful
 - Imagine if your computer froze for hours every time you started a simulation
-

Scheduling Policies

First In-First Out: Process are executed in the order they arrive, and allowed to run to completion

Round Robin: Schedule waiting processes in a circular manner, with each process running for some time slice before being interrupted and rescheduled

Shortest Job First: Of the waiting processes, the one estimated to take the *least amount of time* is allowed to run to completion. Can be preempted if a shorted process arrives.

Priority Levels

- Why should schedulers care about task priority? Isn't this unfair?
 - Some tasks are time sensitive
 - Some tasks are significantly more important than others
 - Interactive applications typically need fast response times to be useful
 - Want to react fast to input
 - CPU bound applications will typically have lower priority
 - Simulations
-

Memory Maps

How do memory maps work?

- What's shared?
 - C Library
 - What's not shared?
 - Heap, stack, Executable code (Most important, executable for lib is shared but the actual executable code running is not)
 - Permissions:
 - R – read, W – write, X – execute
 - Total memory between processes
 - Sum together pages, making sure to only count libraries once over all processes
-

MP3 + etc.

Implementing execute()

1. Generate a PCB
 - PID, parent PID,
 2. Modify paging
 - New user page for program data
 - Flush TLB
 3. Load user program into memory
 4. Modify TSS (esp0)
 - Why don't we touch anything else?
 5. Setup userspace IRET context
 - SS, ESP, EFLAGS, CS, EIP
-



```
ece391> cat frame0.txt
```

```
^/\^/\^/\^/\^/\^/\^/\^/\^/\^/\
      o
    o  o
  o
      o
    o  o
      \
_  _  _  _  _  _  _  _  _  _  _
| \ . \ | \ / / / /
| = _ > \ | \ /
| / \ / | / | /
-----M-----M-----
```

```
ece391> ls_
```



```
ece391> cat frame0.txt
```

```
^/^/^/^/^/^/^/^/^/^/^/^/^/^/^
```

```
o
```

```
o
```

```
o
```

```
o
```

```
o
```


```
  _  
 | \ . \ | \  
 | = > | \  
 | / > | \  
-----
```


```
ece391> ls
```


INT
\$0x80

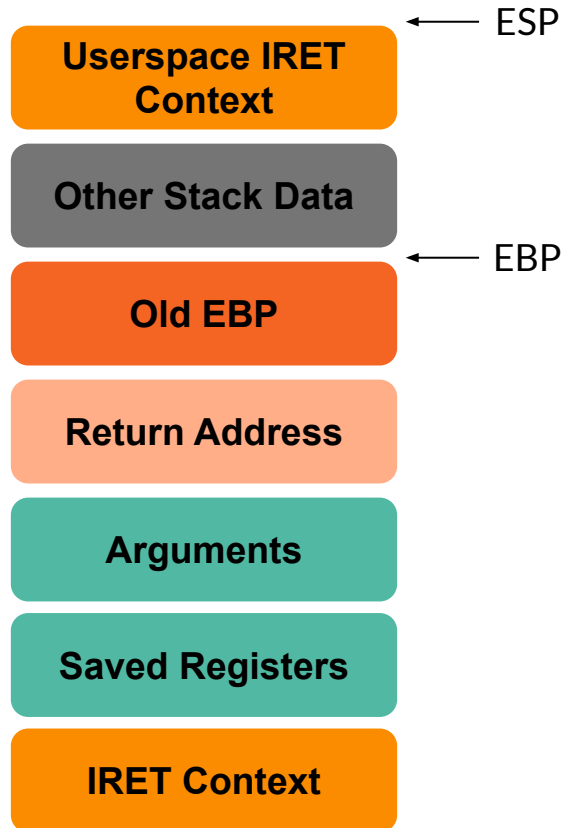
IRET Context


← ESP









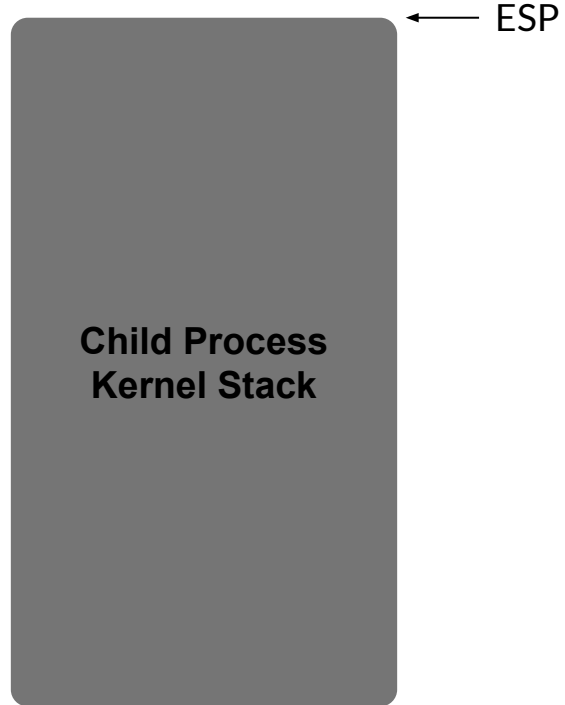


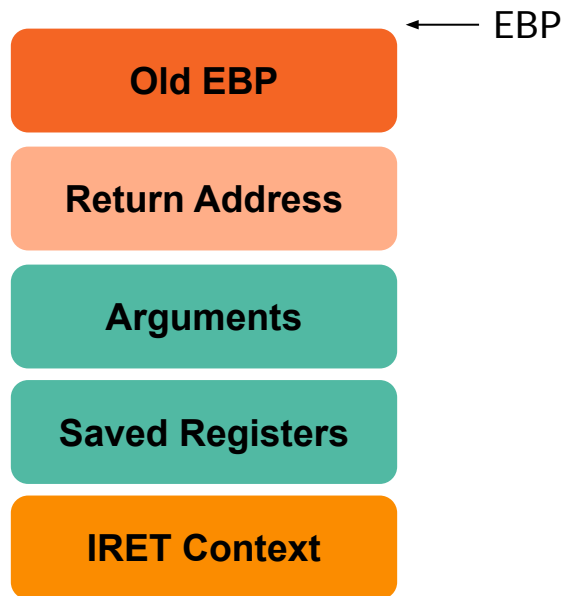


















IRET Context

← ESP



```
ece391> cat frame0.txt
```

```
/\ /\ /\ /\ /\ /\ /\ /\ /\ /\
      o
    o   o
  o
      o   o
    o   o
      \
| \/. \ | \ / / /
| = _ > \ | \ /
| /\ _ / | / | /
-----M-----M-----
```

```
ece391> ls
```

```
.
hello
frame0.txt
...
```

Terminals + Scheduling

- Three terminals -> three leaf processes
 - Up to 6 *total* processes
 - Switch between terminals using keyboard interrupts
 - Remap video memory
 - Round robin scheduling
 - Switch processes on PIT interrupts
 - Takes advantage of IRET context set up by processor
-

Kernel

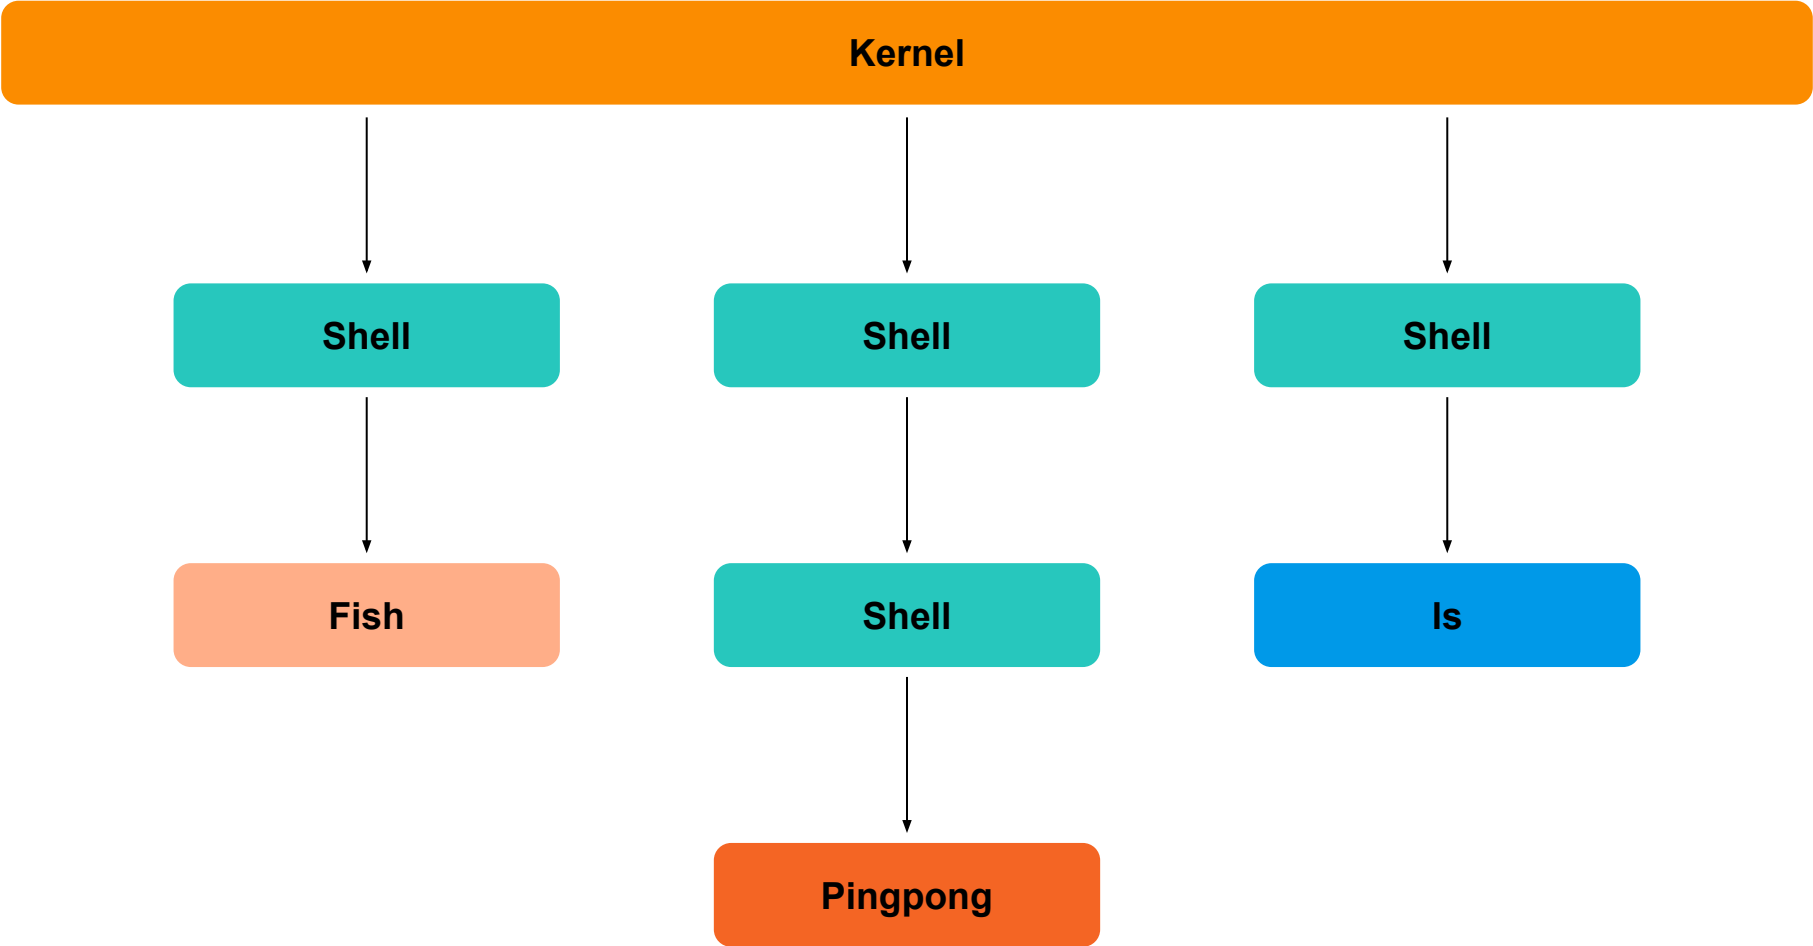
Kernel

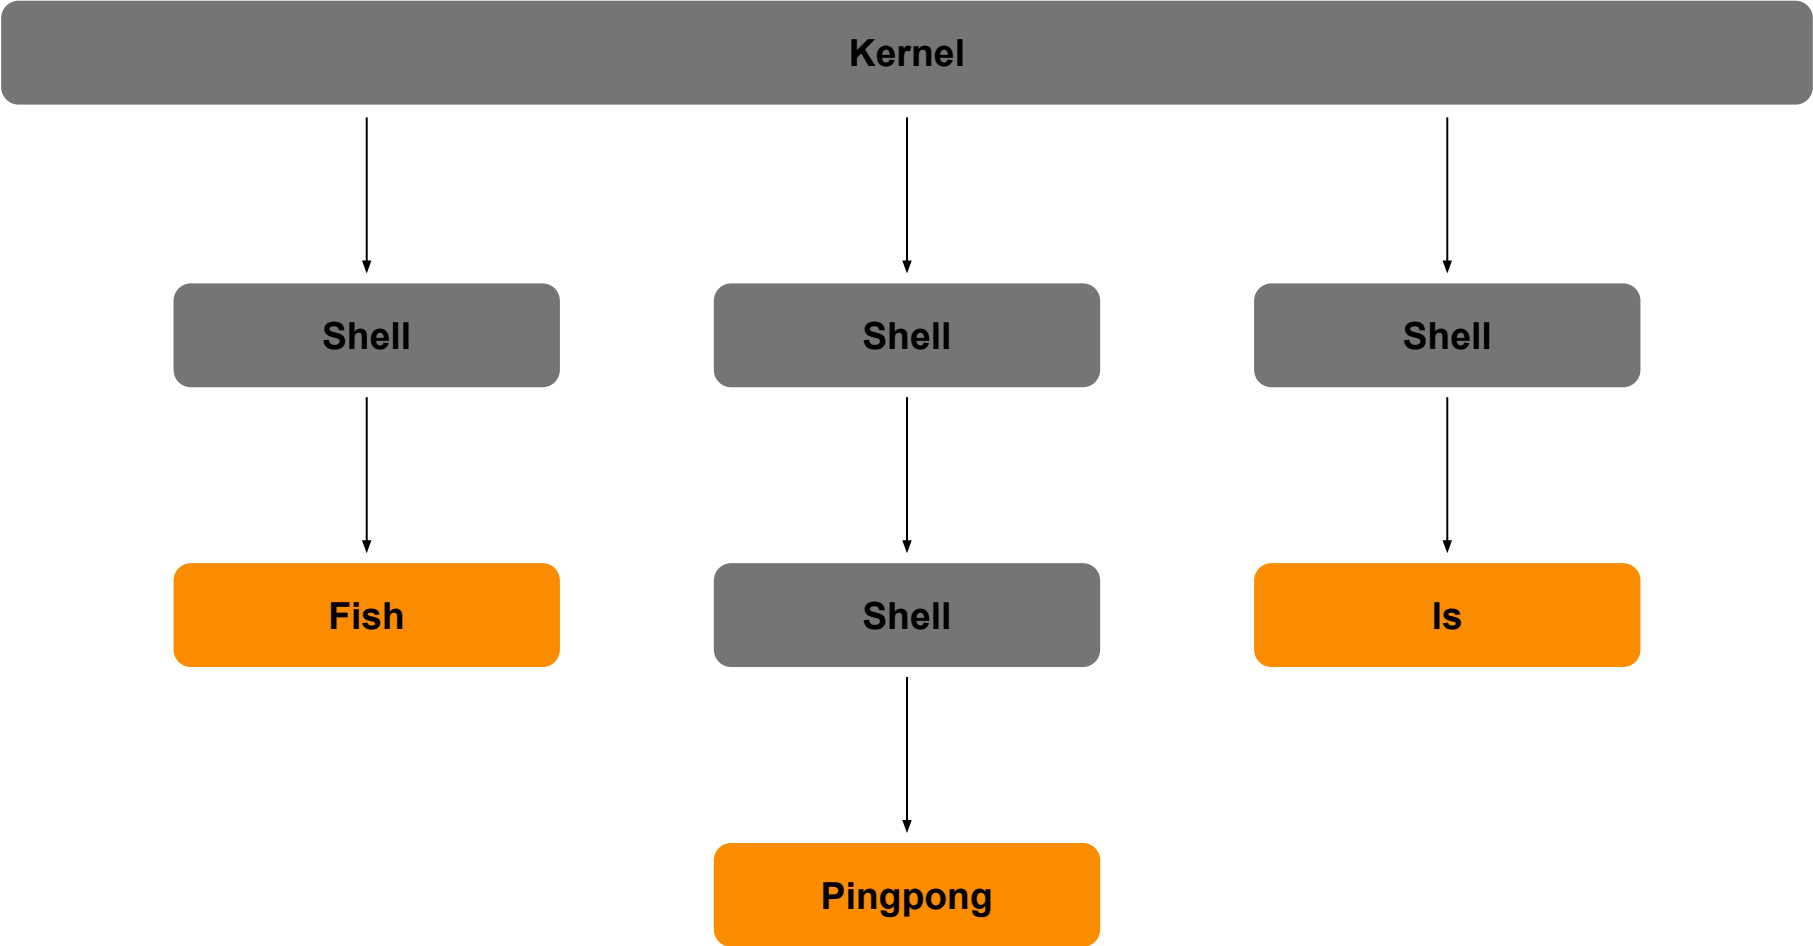
```
graph TD; Kernel[Kernel] --> Shell1[Shell]; Kernel --> Shell2[Shell]; Kernel --> Shell3[Shell];
```

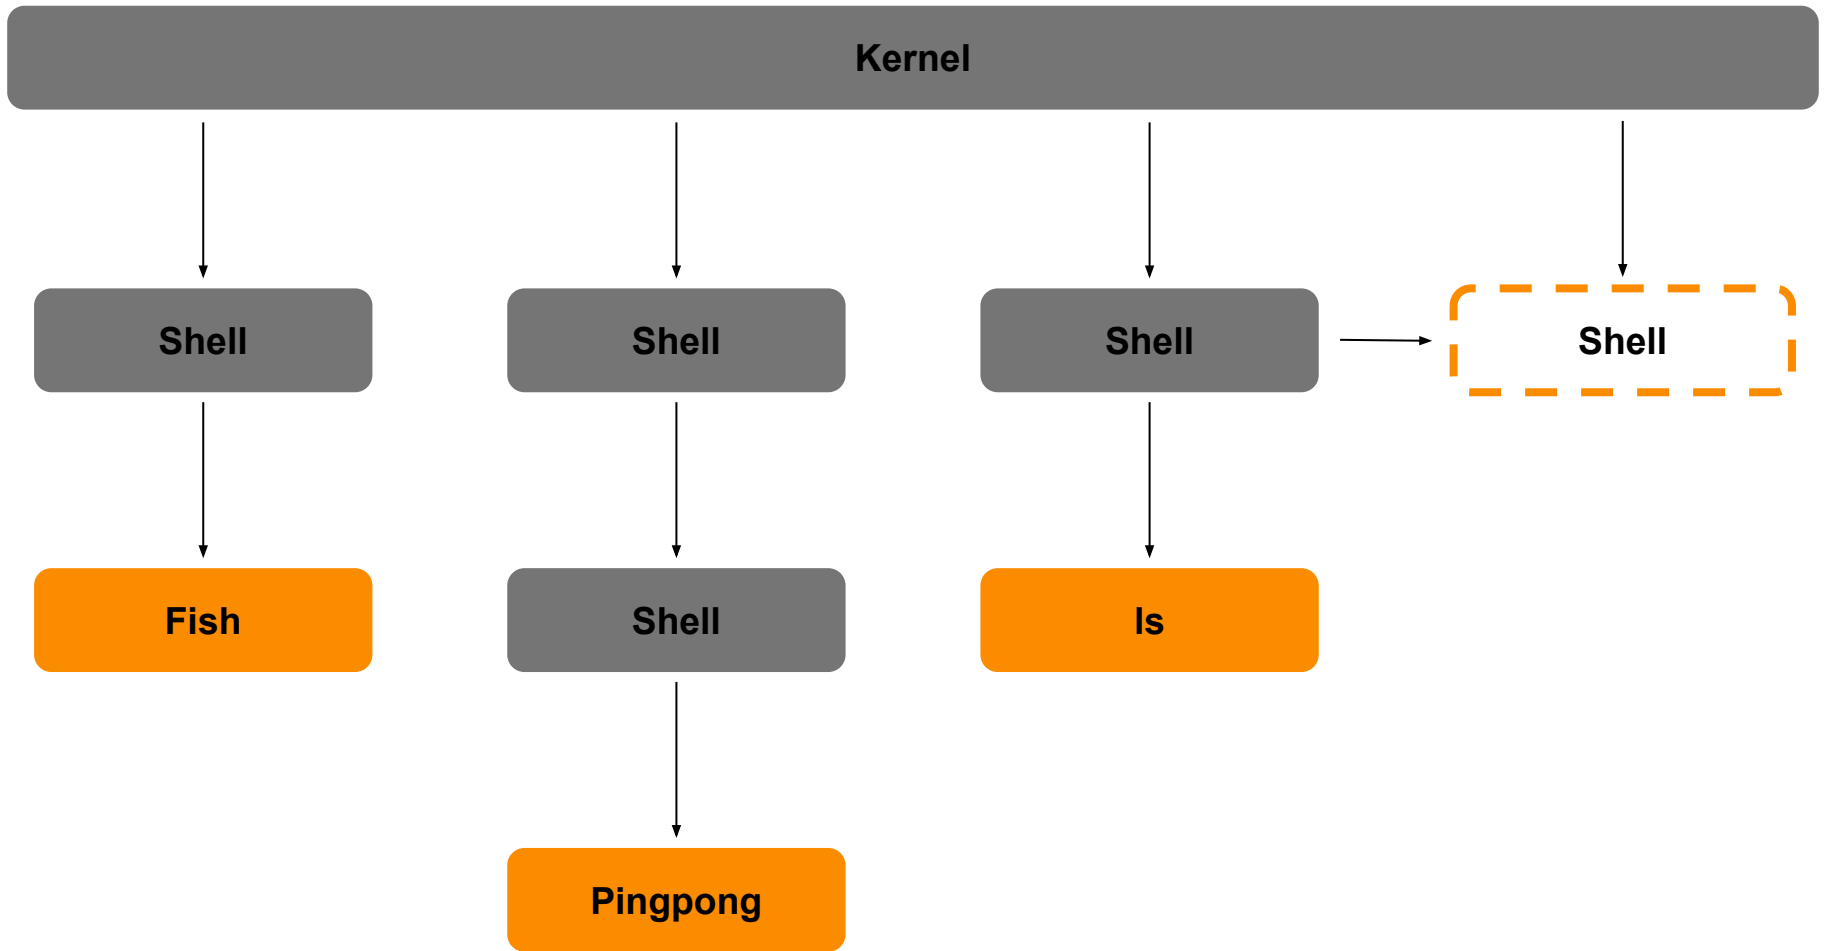
Shell

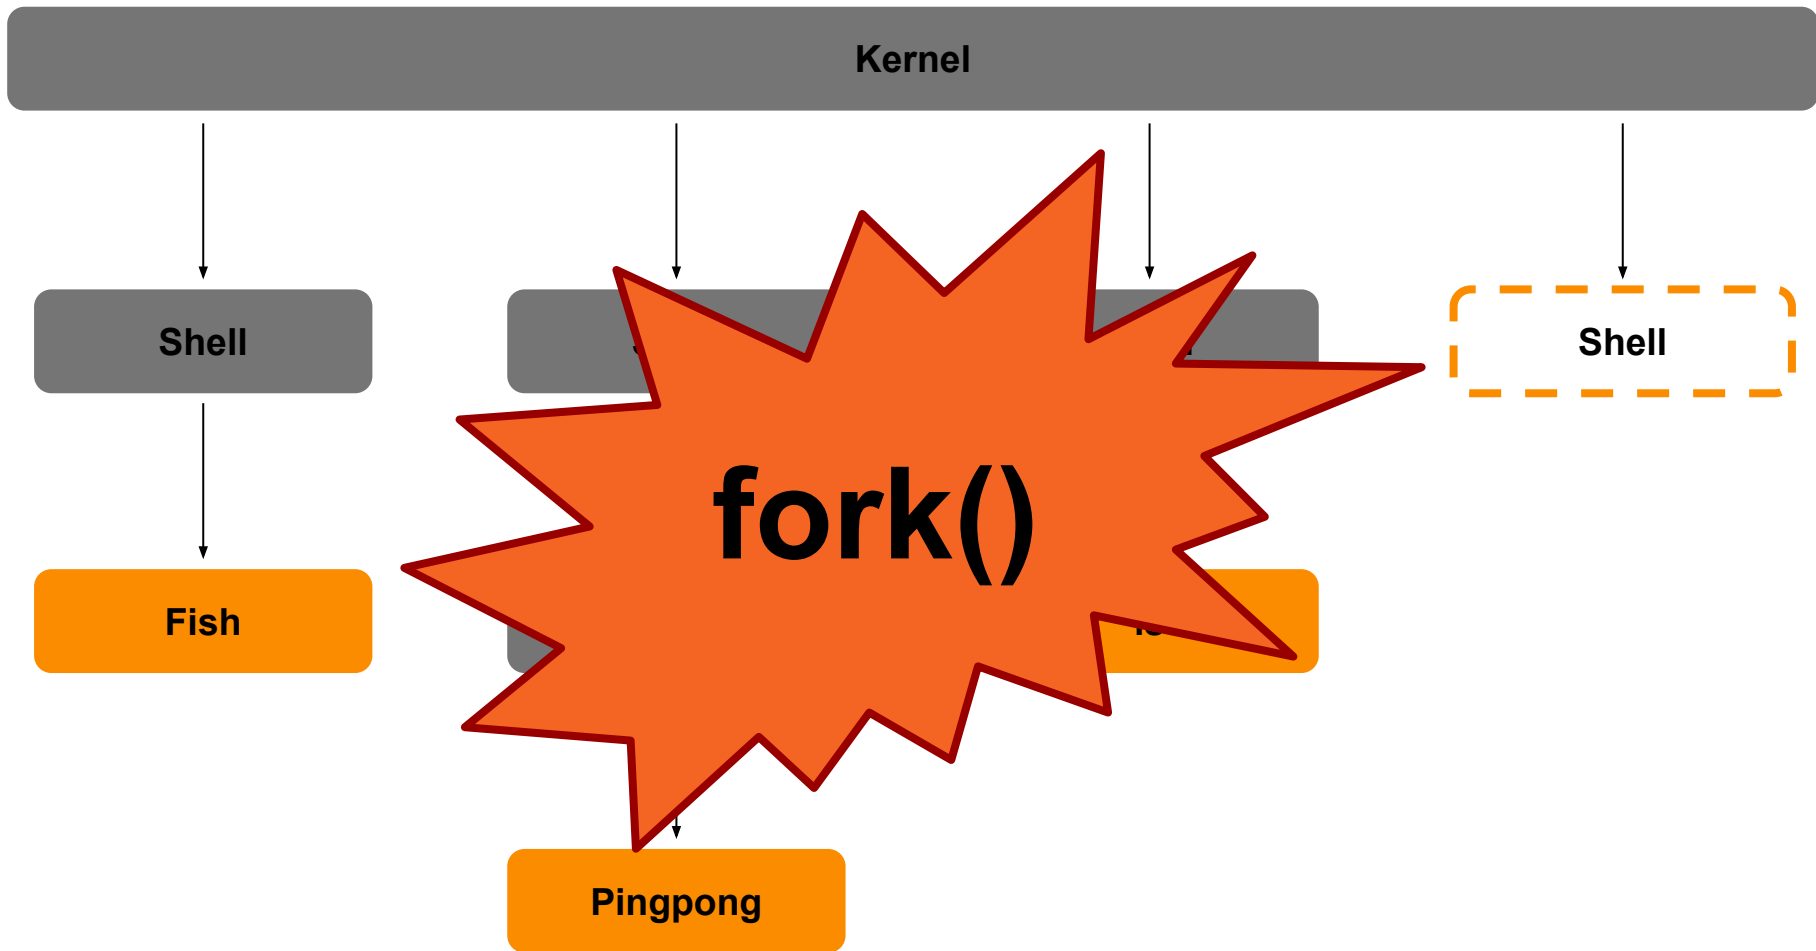
Shell

Shell









The fork() System Call

- Why is fork useful?
 - Create new 'leaf process'
 - Duplicates the process + any metadata associated with it
 - Copy on write
 - Different return value for parent / child
 - Parent gets Child PID, child gets 0
 - Often used in conjunction with exec()
-